

Reliable, Recognised, and Hard to Validate: An Audit of LLM Stance Scores for FOMC Statements

Lenny Lorenz*

October 2026

Abstract

Large language models (LLMs) are now used to score central bank communication. I audit such scores for FOMC statements (2000–2026). Reliability: on a block of 40 statements scored with five prompts, three inexpensive models lose 29% to 45% of their agreement once the rate decision is removed; a frontier model loses 3%. For meeting-to-meeting changes, the choice of model matters about twice as much as the choice of prompt. Recognition: shown statements with dates, names and events removed, the frontier model names the exact month for 98% of them (93% of a paraphrased sample), while a text-similarity baseline recovers at most 43%; it also names the month unprompted in 11.2% of its blinded scoring replies. Blinding and paraphrasing therefore do not rule out look-ahead bias for such documents. Validity: the frontier model’s score changes relate to policy surprises at no-change meetings, but the link fails a multiple-testing adjustment and is concentrated in 2016–2023.

Keywords: large language models, central bank communication, look-ahead bias, measurement reliability, monetary policy surprises.

JEL: E52, E58, C81, G14.

*Email: lennylorenz88@gmail.com. Code, data and every number in this paper are reproducible from the project repository, <https://github.com/lennylrz/fomc-project> (see [paper/README.md](#)).

1 Introduction

Researchers in finance and macroeconomics now use large language models (LLMs) to turn text into numbers: the tone of news headlines, of earnings calls and of central bank statements (Lopez-Lira and Tang, 2023; Hansen and Kazinnik, 2023; Shah et al., 2023). For central bank communication this continues a long line of work that measured statements with word lists, topic models and hand labels (Lucca and Trebbi, 2009; Apel and Blix Grimaldi, 2012; Hansen and McMahon, 2016; Gardner et al., 2022). An LLM promises a stance measure that needs no dictionary and no training data. One asks the model how hawkish a statement is and records the number.

An applied researcher who wants to use that number faces two questions. The first is about measurement: would another, equally reasonable prompt or model give the same answer? The second is about what the model actually does: did it read the text, or did it recognise the document and remember what happened next? The second question is the look-ahead problem that Glasserman and Lin (2024), Sarkar and Vafa (2024) and Ludwig et al. (2026) raise for any LLM measure applied to documents written before the model’s training data end. The usual remedies either change the input after the fact, by removing names and dates or by paraphrasing, or use models whose training data stop before the documents (He et al., 2025; Drinkall et al., 2024).

This paper audits LLM stance scores for the post-meeting statements of the Federal Open Market Committee (FOMC) with these two questions in mind. I score 221 policy statements from 2000 to 2026 with five prompts, using a frontier model (Claude Fable 5.1, “Fable” below), that is, a current top-tier general-purpose model, and an inexpensive one (Gemini 2.5 Flash). On a block of 40 statements from 2015–2019, two more inexpensive models (Claude 3 Haiku and GPT-4o mini) score every statement with every prompt as well. GPT-4, which has a documented training cutoff of September 2021, scores the full corpus with one prompt, and a word-count dictionary serves as a baseline that cannot remember anything. I then ask three things: how reliable the scores are, whether the frontier model recognises the statements it scores, and whether its scores relate to market reactions. The paper has three findings.

Reliability depends on the model. I measure agreement across prompts with the intraclass correlation (ICC), which is 1 when every prompt gives the same scores up to a constant and falls towards 0 as prompt-specific noise grows. On a block of 40 statements from 2015–2019, the three inexpensive models agree moderately on score levels (ICC between 0.609 and 0.810), but much of that agreement comes from the rate decision itself. Once the decision is regressed out, their ICC falls by 29% to 45%. Fable’s falls by 3%. For meeting-to-meeting changes in the score, a three-way variance decomposition attributes 15.2% of the variance to models disagreeing about which statements became more hawkish and 6.7% to prompts disagreeing. Model choice matters about twice as much as prompt choice. How much agreement survives depends on the sample: on the full corpus Gemini’s drop is only 6%.

Recognition survives the standard de-contamination tools. This is the main new result. I remove month names, years, member names and named events from each statement (“blinding”) and ask Fable when the meeting took place. It names the exact month for 202 of 206 pre-2024 statements, including 55 of the 56 statements from the zero-lower-bound years, when the policy

rate was the same at every meeting. After an LLM paraphrase that changes the wording and states the rate only as a change, it still names 28 of 30. A text-similarity baseline with no memory places most statements within three months but can recover the exact meeting for at most 43% of them. Fable also recognises statements without being asked: in 11.2% of its scoring replies on blinded text it volunteers the statement’s month and year (“This is the January 2014 FOMC statement”), and its replies on the original text sometimes cite information from outside the statement, such as the decisions that followed. Its knowledge extends to roughly early 2026, so even the 2024–2025 statements are not “after its cutoff”. The lesson reaches beyond this corpus, although it is shown here for one model: for well-known documents, changing the text after the fact does not rule out look-ahead bias.

Validity is weak and cannot be separated from contamination. Fable’s meeting-to-meeting score change (dScore) is positively related to the monetary policy surprise (MPS) of Bauer and Swanson (2023b), which measures how far interest-rate futures move in the 30 minutes around an FOMC announcement, at meetings where the rate did not change. In the regression designated as primary, the t -statistic with autocorrelation-robust (Newey–West) standard errors is 2.39 ($p = 0.018$, $n = 132$). This is suggestive, not established. Adjusted with Holm’s method for the twelve regressions in its family, the p -value is 0.199; the specification was frozen only after earlier specifications had been seen; and the relation is concentrated in 2016–2023 (t of 3.40, against 1.32 for 2000–2007 and 0.42 for 2008–2015). Daily changes in 2- and 10-year Treasury yields are not significantly related to dScore under the primary corpus definition (t of 1.03 and 1.74); they become significant if two borderline unscheduled statements are dropped, or if non-policy notices are kept. Fable’s scores track market surprises better than Gemini’s (in an exploratory comparison the difference in standardised slopes has z between 2.23 and 2.59), but Fable agrees no better than Gemini or GPT-4 with published human labels of the same statements (Shah et al., 2023). That pattern fits a model that knows more about what happened at least as well as a model that reads better.

What the design cannot do. The paper does not show that Fable’s market link comes from memory, and it does not show that it does not. Recognition is necessary for recalling a market reaction, but it is not evidence of such recall. The fall in the market link on paraphrased text mixes two effects: the removal of exact wording, which markets do respond to, and the removal of cues to the statement’s identity. GPT-4, the one model with a documented cutoff, answers on so coarse a scale that it cannot serve as a control, and the 2024–2026 statements are both too few (power of about 8% to 11%) and, for Fable, not unseen. Only models trained on data that stop before the documents, or truly unseen documents, can settle the question.

1.1 Related literature

LLMs as measurement instruments. Lopez-Lira and Tang (2023) show that ChatGPT’s reading of news headlines predicts next-day stock returns, and Hansen and Kazinnik (2023) find that GPT models classify the stance of FOMC communication better than standard text methods. Shah et al. (2023) hand-label sentences from FOMC documents as hawkish, dovish or neutral and build a stance measure from a fine-tuned model. Ludwig et al. (2026) set out when LLM outputs can be used in econometric work: prediction tasks need the absence of training leakage, while measurement

tasks need a validation sample. This paper takes the measurement view and asks how much the measure depends on arbitrary choices. The intraclass correlation and Krippendorff’s alpha, two standard statistics of agreement between raters, are the tools for that question (Shrout and Fleiss, 1979; Hayes and Krippendorff, 2007; Krippendorff, 2018).

Look-ahead bias. Glasserman and Lin (2024) separate look-ahead bias from a “distraction” effect in GPT sentiment scores of news headlines and use anonymised company names as a remedy. Sarkar and Vafa (2024) propose tests for look-ahead bias in pretrained models. He et al. (2025) and Drinkall et al. (2024) train models only on text available at each date, so that look-ahead bias is ruled out by construction. Work on data contamination in machine learning shows that models can reproduce training text and recognise benchmark items (Carlini et al., 2021; Golchin and Surdeanu, 2024). My contribution here is a direct test of the anonymisation remedy on a corpus of well-known documents: for such documents, removing identifiers does not stop identification, and identification also shows up, unprompted, in the scoring replies themselves.

Central bank communication and market surprises. Statements move asset prices beyond the rate decision (Gürkaynak et al., 2005), and FOMC surprises can be split into rate, forward-guidance and asset-purchase components (Swanson, 2021). High-frequency surprises may also carry information about the economy rather than about policy (Nakamura and Steinsson, 2018; Jarociński and Karadi, 2020); Bauer and Swanson (2023a) argue instead that the Fed and markets respond to the same public news. I use the surprise series of Bauer and Swanson (2023b), including their version purged of predictable components, as the main outcome. Text measures of Fed communication have been related to markets and the macroeconomy before (Lucca and Trebbi, 2009; Hansen and McMahon, 2016; Gardner et al., 2022). Relative to that work, the question here is not whether statements contain information, but whether an LLM score of them can be validated against markets when the model may have seen the market reaction.

The rest of the paper is organised as follows. Section 2 describes the statements and the market data, and Section 3 the scoring design. Section 4 reports reliability, Section 5 the recognition audit, and Section 6 the market validation. Section 7 discusses what this design can and cannot test, and Section 8 concludes. The appendix gives the prompts verbatim, the corpus classification, the parse audit, additional results and a full list of disclosures.

2 Data

2.1 FOMC statements

The corpus holds 228 statements released by the FOMC between 2000-02-02 and 2026-07-29, taken from the press-release pages of the Federal Reserve Board. I classified each by hand after reading it (Appendix Table A1): 212 are post-meeting statements of scheduled meetings, 7 announce unscheduled changes in the federal funds rate, 2 announce unscheduled policy decisions without a rate change (the risk assessment of 17 August 2007 and the open-ended asset purchases of 23 March 2020), and 7 are notices about liquidity facilities, swap lines, technical operations or the policy framework. The *policy corpus* drops the notices and keeps 221 statements. It is used for all reliability statistics and market regressions; the recognition counts in Section 5 cover every file, notices included. Section 6.4 shows results for a *strict* corpus that also drops the two unscheduled

policy decisions (219 statements) and for all 228 files.

Two corrections were made to the raw corpus. First, the Board’s page for the statement of 27–28 June 2007 carries the date 18 June in its address, so the statement was first filed under the wrong day; it is dated 28 June here, which also gives it its market data. Second, the text of 51 statements from 2000–2005 originally ended with a line of page navigation that included a year (in 50 cases the statement’s own year). That line was removed, but only after the five-prompt scoring runs had been made, so those runs saw it. The single-prompt re-scores, the GPT-4 run and the dictionary use the cleaned text (Table 2).

The texts of the statements of 17 June and 29 July 2026 are much shorter than earlier ones and could not be checked against the Board’s website from the computing environment used here. They read as complete statements in a new, shorter format. They enter the post-2024 checks (Section 7), the dating probe on recent statements (Section 5.4) and the comparison of prompt spread on recent statements (Section 4); the last is reported with and without them.

2.2 Market data

The main outcome is the monetary policy surprise (MPS) of Bauer and Swanson (2023b). MPS summarises the change in short-term interest-rate futures in a 30-minute window around each FOMC announcement, in percentage points; a positive value means markets read the announcement as tighter than expected. Bauer and Swanson (2023b) show that MPS is partly predictable from data available before the announcement and also provide MPS_ORTH , the part of MPS orthogonal to those data. The same dataset gives the 30-minute changes in the 2- and 10-year Treasury yields (t_2 , t_{10} , in basis points); their sign and scale match the daily yield changes, with which they correlate at 0.64 (t_2 , d_2) and 0.67 (t_{10} , d_{10}) on statement days. These series end in December 2023.

Daily yields are the constant-maturity 2- and 10-year Treasury yields from the Federal Reserve’s H.15 release (FRED series DGS2 and DGS10). The daily outcomes d_2 and d_{10} are close-to-close changes in basis points on the statement day, from the previous trading day. The yield file matches the FRED series to the hundredth on all 1229 days on which they overlap (2021-10-06 to 2026-09-08). A statement released on a non-trading day (15 March 2020, a Sunday) takes the next trading day’s change. The daily series used here starts in 2003, so the daily regressions have fewer meetings than those using MPS.

A *hold meeting* is a statement at which the federal funds rate target did not change. Hold meetings matter for this paper because at them the stance score cannot simply echo a rate move. They are not free of policy news: a hold can itself surprise markets, and the hold sample includes announcements of asset purchases (for example 18 March 2009) and the unscheduled statements of 17 August 2007 and 23 March 2020. Before 2024 the policy corpus contains 200 statements, which give 199 meeting-to-meeting changes; 132 of these are hold meetings with an MPS. Table 1 describes the outcomes.

Table 1: Market outcomes, policy corpus, statements to December 2023

Outcome	Unit	Sample	n	Mean	SD	Min	Max	Years
MPS	pp	all	198	-0.002	0.059	-0.283	0.152	2000–2023
		hold	132	0.002	0.039	-0.111	0.152	2000–2023
MPS_ORTH	pp	all	189	0.003	0.052	-0.232	0.161	2000–2023
		hold	125	0.003	0.041	-0.107	0.161	2000–2023
2-year yield, 30-minute (t_2)	bp	all	198	-0.8	5.0	-16.9	20.5	2000–2023
		hold	132	-0.7	4.4	-13.7	20.5	2000–2023
10-year yield, 30-minute (t_{10})	bp	all	198	-0.3	4.2	-29.6	16.4	2000–2023
		hold	132	-0.4	4.4	-29.6	16.4	2000–2023
2-year yield, daily (d_2)	bp	all	173	-0.9	7.7	-28.0	23.0	2003–2023
		hold	120	-0.8	6.1	-27.0	17.0	2003–2023
10-year yield, daily (d_{10})	bp	all	173	-0.9	8.0	-51.0	22.0	2003–2023
		hold	120	-0.5	8.0	-51.0	20.0	2003–2023

Notes: “all” is every statement with a predecessor in the corpus; “hold” is the subset at which the federal funds rate target did not change. MPS and MPS_ORTH are from Bauer and Swanson (2023b), in percentage points. t_2 and t_{10} are 30-minute changes in the 2- and 10-year Treasury yields around the announcement, from the same dataset; d_2 and d_{10} are daily close-to-close changes of the H.15 constant-maturity yields (FRED DGS2, DGS10).

3 Scoring design

3.1 Models and prompts

Every statement is scored by asking a model for a single number between -1 (dovish) and $+1$ (hawkish). I use five prompts that ask the same question in different words (Appendix A gives them verbatim). P1 asks for the score with no further instruction. P2 first defines hawkish and dovish. P3 asks the model to answer as a fixed-income analyst. P4 reverses the scale, so that $+1$ means dovish; its answer is flipped back in code. P5 asks the model to list the hawkish and dovish elements before giving the score. All calls use temperature 0 through the OpenRouter interface.

Table 2 lists the scoring runs. The reliability block crosses 40 policy statements from 2015–2019 with all five prompts and four models: Gemini 2.5 Flash, Claude 3 Haiku, GPT-4o mini and Claude Fable 5.1. Gemini and Fable also score the full corpus. Gemini has all five prompts for every policy statement. Fable has all five for 174 statements and between one and four for the other 47, all of them from 2020 onwards. GPT-4, listed by OpenRouter with a training cutoff of 30 September 2021, scores every statement with P1. Fable then re-scores the pre-2024 statements with P1 three times: on the cleaned original text, on a blinded version and on a paraphrased version (Section 3.3).

Replies are parsed strictly: the reply must end in a number in $[-1, 1]$, or contain a line “SCORE: x ”. A reply that is cut off, ends in a bare “0.” or gives a value outside the range is treated as missing. Until 3 October 2026 the scoring scripts stored only the first 300 characters of each reply, so longer replies could not be re-parsed afterwards. Re-parsing every shorter stored reply and reading one longer reply by hand removed 5 cells (Appendix Table A2).

3.2 From scores to dScore

For each statement and prompt I average repeated runs of the same cell; Fable’s test–retest agreement on the 245 cells scored twice is high (correlation 0.996, mean absolute difference 0.024).

Table 2: Scoring runs

Model (OpenRouter ID)	Use	Prompts	Statements	Text
google/gemini-2.5-flash	Reliability block	5	41	original
anthropic/claude-3-haiku	Reliability block	5	41	original
openai/gpt-4o-mini	Reliability block	5	41	original
anthropic/claude-fable-5.1	Reliability block	5	41	original
google/gemini-2.5-flash	Reliability, 2004–07	5	34	original
google/gemini-2.5-flash	Full corpus	5	228	original
anthropic/claude-fable-5.1	Full corpus	5	228	original
openai/gpt-4	Full corpus (cutoff model)	1	228	original, cleaned
anthropic/claude-fable-5.1	Re-score, pre-2024	1	206	original, cleaned
anthropic/claude-fable-5.1	Re-score, pre-2024	1	206	blinded
anthropic/claude-fable-5.1	Re-score, pre-2024	1	206	paraphrased
word count (no model)	Full corpus (no LLM)	1	228	original, cleaned

Notes: Model IDs as listed by OpenRouter. Statements = distinct statement files with a score (all 228 files, including the notices that the policy corpus drops). The full-corpus Fable file has all five prompts for 174 policy statements and one to four for 47. “Original” text is the text as first downloaded; for 51 statements from 2000–2005 it ended with a navigation line that included a year. “Cleaned” text has that line removed. Blinded and paraphrased versions are described in Section 3.3. The dictionary is a word-count score (Section 3.4).

The statement’s score is the mean over its prompts after removing each prompt’s average offset, estimated on the statements that have all five prompts (prompt fixed effects). Without this step, a statement scored with a different set of prompts than its predecessor would show a spurious change. The score change, $dScore$, is the score of a statement minus the score of the previous statement in the policy corpus. Notices are removed before the differences are taken, so each statement is compared with the previous policy statement. A positive $dScore$ means the statement reads as more hawkish than its predecessor. Section 6.4 shows that the main results do not depend on the prompt fixed effects.

3.3 Text versions

The *blinded* version of each statement removes the cues that most directly date it. Month names become “[month]”; years become offsets from the statement’s own year (“[year+1]”), so that date-based guidance keeps its meaning; the names of all FOMC members become “[member]”; the roster of members voting in favour is dropped, while dissents are kept without names; and named one-off events become “[event]”. Rate levels, purchase amounts and the phrasing are left as they are.

The *paraphrased* version rewrites each blinded statement with Gemini 2.5 Flash. The instructions (Appendix A) ask for the same economic content in a new order and new words, with no reused phrase of four or more words, and with rate levels and asset holdings stated only as changes from the previous meeting. To keep the decision right, the prompt states the true change in the federal funds target. Automatic checks flag rewrites that are too long, add dissents or a 2 percent objective that the original lacks, or keep rate levels; a flagged rewrite is tried once more. Of the 206 paraphrases, 204 pass the checks and 2 (2013-07-31 and 2019-10-11) still keep a number after the retry: a guidance threshold in the first case, an operational rate in the second, which is a notice outside the policy corpus. A retry was needed for 3 rewrites. Their four-word overlap with the blinded text is 0.120 on average (maximum 0.395), and they are 0.79 times as long. In all, 11 paraphrases keep the 6.5 percent unemployment-rate threshold from the forward guidance of 2012–2014; the checks did not

treat a threshold as a level. I read a sample of paraphrases against their originals, including those on which an earlier attempt had gone wrong, and found the content faithful.

3.4 Baselines

Two baselines involve no LLM. The *dictionary* score counts hawkish and dovish word combinations (for example, “inflation” near “rise” counts as hawkish) in the spirit of Apel and Blix Grimaldi (2012), using a short word list written for this project, and scales the net count to $[-1, 1]$. It cannot have memorised anything. The *text-similarity* dating baseline, used in Section 5, gives each blinded statement the date of the most similar other blinded statement. Finally, Shah et al. (2023) hand-label sentences from FOMC minutes, speeches and press conferences as hawkish, dovish or neutral. Because the minutes repeat the statement, 423 of their labelled sentences appear word for word in 152 of the statements here. This gives a human benchmark for score levels that neither markets nor model memory can affect.

4 Reliability across prompts and models

4.1 Measures

I measure agreement across the five prompts in two ways. The intraclass correlation ICC(3,1) of Shrout and Fleiss (1979) treats the prompts as fixed raters and asks what share of the variance in a single prompt’s scores is variance between statements, after removing each prompt’s average level. It equals 1 when all prompts rank and space the statements identically, up to a constant shift. Krippendorff’s alpha with the interval metric (Hayes and Krippendorff, 2007; Krippendorff, 2018) also counts a constant shift between prompts as disagreement, so it is at or below the ICC when prompts differ in level.

Each measure is computed four ways: on raw score levels; on levels after the rate decision is regressed out (“resid”); on hold meetings only; and on meeting-to-meeting changes. For “resid” each prompt’s scores are regressed on a single code for the decision (+1 hike, 0 hold, −1 cut). Separate hike and cut dummies give similar results (below). The *drop* is $1 - \text{resid}/\text{raw}$. It asks how much of the agreement is agreement about the decision itself, which no one needs an LLM to read.

4.2 Results

Table 3 and Figure 1 report the results. The first block crosses the same 40 policy statements from 2015–2019 with all five prompts and four models. On raw levels the three inexpensive models reach an ICC between 0.609 and 0.810. Once the rate decision is removed, their ICC falls to between 0.333 and 0.562, a drop of 29% to 45%. Fable starts at 0.957 and keeps 0.929 (a drop of 3%). The block contains 9 hikes and 3 cuts. With separate hike and cut dummies instead of the single code, the drops are 26% to 39% for the inexpensive models and 3% for Fable. Alpha tells the same story. Its drops for the inexpensive models are smaller (16% to 29%) because alpha already penalises their level differences in the raw scores, but on hold meetings it shows how little some of them agree: alpha is 0.084 for Claude 3 Haiku and 0.355 for GPT-4o mini, against 0.928 for Fable. On

Table 3: Agreement across five prompts

Model	Window	n	ICC(3,1)					Krippendorff's α				
			Raw	Resid	Hold	Chg	Drop	Raw	Resid	Hold	Chg	Drop
<i>Crossed block: same statements, five prompts, four models</i>												
Gemini 2.5 Flash	2015-01–2019-12	40	0.810	0.562	0.580	0.697	31%	0.791	0.564	0.524	0.698	29%
Claude 3 Haiku	2015-01–2019-12	40	0.609	0.333	0.420	0.419	45%	0.398	0.336	0.084	0.422	16%
GPT-4o mini	2015-01–2019-12	40	0.728	0.517	0.537	0.577	29%	0.634	0.519	0.355	0.579	18%
Claude Fable 5.1	2015-01–2019-12	40	0.957	0.929	0.937	0.920	3%	0.953	0.930	0.928	0.921	2%
Gemini 2.5 Flash	2004-01–2007-12	33	0.912	0.830	0.907	0.696	9%	0.891	0.831	0.885	0.697	7%
<i>Full corpus (statements with all five prompts)</i>												
Gemini 2.5 Flash	2000-02–2026-07	221	0.934	0.874	0.885	0.729	6%	0.933	0.874	0.881	0.729	6%
Claude Fable 5.1	2000-02–2026-07	174	0.976	0.957	0.961	0.896	2%	0.972	0.957	0.950	0.896	2%

Notes: Policy corpus; statements with all five prompts. Raw = score levels; Resid = levels after regressing each prompt's scores on the rate decision coded +1/0/−1; Hold = hold meetings only; Chg = changes between consecutive corpus statements; Drop = 1 − Resid/Raw. ICC(3,1) follows Shrout and Fleiss (1979); alpha uses the interval metric.




Figure 1: Agreement across prompts by model, 2015–2019 block ($n = 40$ statements). Part of the inexpensive models' agreement on levels is agreement on the rate decision.

meeting-to-meeting changes, the quantity that market regressions use, the inexpensive models reach an ICC of 0.419 to 0.697 and Fable 0.920.

The size of the drop depends on the sample. For Gemini it is 31% in the 2015–2019 block but 9% on 2004–2007 and 6% on the full corpus. ICC is a ratio of between-statement to total variance, so it is high when the statements in the sample differ a lot in stance, as they do across the whole period from 2000 to 2026. Within a few years of similar statements, prompt noise weighs more. The reliability of an LLM score is therefore a property of the model, the prompt set and the sample together, and a researcher should report it for the sample actually used.

Fable's consistency is also high in a second sense: the 245 cells it scored twice in separate runs correlate at 0.996.

Table 4: Variance decomposition: statement $\times$ prompt $\times$ model, 2015–2019 block

Data	Scaling	n	Statement	Prompt	Model	S $\times$ P	S $\times$ M	P $\times$ M	Residual
Levels	raw	40	61.0	4.5	0.0	3.0	9.7	5.7	16.1
Levels	std. within model	40	61.3	4.6	0.0	3.0	9.2	6.0	15.9
Changes	raw	39	49.1	0.0	0.0	6.7	15.2	0.0	29.0
Changes	std. within model	39	50.2	0.0	0.0	6.4	15.2	0.0	28.3

Notes: Shares of total variance (%), random-effects ANOVA, 40 statements $\times$ 5 prompts $\times$ 4 models (Gemini 2.5 Flash, Claude 3 Haiku, GPT-4o mini, Claude Fable 5.1); negative variance estimates are set to zero. “Std. within model” z-scores each model’s scores first, which removes scale differences between models. S $\times$ P: prompts disagree about which statements are more hawkish; S $\times$ M: models disagree; the residual contains the three-way interaction and noise.

4.3 Model choice versus prompt choice

Prompt sensitivity is only one source of arbitrariness. A researcher also chooses a model. On the crossed block I decompose the variance of the scores into statement, prompt and model effects and their interactions, with a random-effects analysis of variance. With one score per cell, the three-way interaction cannot be separated from noise, so it is part of the residual. Table 4 reports the shares.

On levels, agreed differences between statements account for 61.0% of the variance. On changes, the agreed part falls to 49.1%. The main disagreement on changes is between models: the statement-by-model interaction, which measures models disagreeing about which statements became more hawkish, takes 15.2%, against 6.7% for the statement-by-prompt interaction. The residual rises from 16.1% on levels to 29.0% on changes. Standardising each model’s scores first leaves these shares almost unchanged, so they do not reflect differences in scale. For meeting-to-meeting changes, the choice of model matters about twice as much as the choice of prompt. A robustness check across prompts alone understates how much an LLM stance measure depends on the researcher’s choices.

4.4 Is the frontier model consistent because it recognises the text?

One reading of Fable’s high reliability is that it is simply a better reader. Another is that it recognises each statement and recalls a fixed view of it, whatever the prompt (Section 5 shows it does recognise them). The two readings predict different behaviour on statements the model has not seen. Fable’s dating probe (Section 5.4) suggests that its knowledge fades from March 2026, so I scored the seven latest statements with all five prompts and split them by whether the probe dated them exactly and with high confidence. The prompt spread (the standard deviation of the five scores) averages 0.085 on the 4 statements it probably has not seen (March to July 2026) and 0.094 on the 3 it dates confidently, against a median of 0.040 and a 90th percentile of 0.117 on the 171 pre-2024 statements with all prompts (Appendix Table A9). All seven recent statements lie between the 0.44 and 0.94 quantiles of the pre-2024 spread, so Fable is somewhat less consistent on recent statements than on old ones, whether it knows them or not. Two of the four statements it probably has not seen are the unverified texts of June and July 2026 (Section 2); without them, the mean spread of the remaining 2 is 0.117, slightly above that of the known statements. Statements this few cannot settle the question.

Table 5: Dating blinded and paraphrased statements

	Claude Fable 5.1, asked for the date			Text-similarity
	Blinded 2000–2023	Paraphrased 2000–2023 (sample)	Blinded 2024–2026	baseline 2000–2023
Exact meeting	202/206 (98.1%)	28/30 (93.3%)	19/22 (86.4%)	$\leq$ 89/206 (43.2%) ^a
Within 1 month	204/206 (99.0%)	29/30 (96.7%)	21/22 (95.5%)	106/206 (51.5%)
Within 3 months	206/206 (100.0%)	30/30 (100.0%)	22/22 (100.0%)	190/206 (92.2%)
ZLB statements, exact	55/56 (98.2%)	6/8 (75.0%)	–	$\leq$ 28/56 (50.0%) ^a
Confidence “high”	205/206 (99.5%)	29/30 (96.7%)	17/22 (77.3%)	–

Notes: Fable is asked for the year and month of the meeting. Paraphrased = an evenly spaced sample of 30 paraphrased statements. 2024–2026 = all blinded statements from 2024 on. ZLB (zero-lower-bound) statements: January 2009 to October 2015, when the target range was the same at every meeting (1 of them a notice). The 2000–2023 columns include 6 notices. Text-similarity baseline: each blinded statement takes the date of its most similar other blinded statement (TF-IDF over words and word pairs, cosine). ^aStatements whose two nearest neighbours are the meetings just before and after, an upper bound on exact recovery for this method.

5 Recognition audit

A model that recognises a statement could, in principle, score it with knowledge of what happened next: the market reaction, the commentary, the following decisions. The standard response is to hide the statement’s identity before scoring. This section tests whether that works for FOMC statements. It asks two questions: can Fable identify a blinded or paraphrased statement when asked, and does it identify statements without being asked while it scores them?

5.1 Dating on request

Fable is shown each blinded pre-2024 statement and told that months, years, names and events have been replaced by placeholders. It is asked to identify the meeting, to give its best guess even if unsure, and to end with the year, the month and a confidence level (Appendix A). The probe covers all 206 pre-2024 files, including 6 notices. The exact month identifies the meeting except in the 7 months that have two statements. The same probe is run on a sample of 30 paraphrased statements spread evenly over the period, with a note that the text may have been reworded.

A model with no memory can still date statements roughly. Policy-rate levels, purchase amounts and the boilerplate of each era change slowly, so neighbouring statements look alike. To measure how far the text alone goes, the text-similarity baseline represents each blinded statement by the frequencies of its words and word pairs (TF-IDF weights), finds the most similar *other* blinded statement by cosine similarity, and takes that statement’s date. The baseline is generous: it sees the true date of every other statement. It can never name the exact meeting, because its answer is always another meeting, so I count the statements whose two nearest neighbours are the meetings just before and just after. For those, a reader who knew the sequence of statements could in principle infer the exact meeting. This count is an upper bound for this method, not for every reader without memory: someone who knows the history of the policy rate, the purchase programmes and the guidance dates could date many statements more precisely without having read them before. The baseline shows how far textual similarity alone goes.

Table 5 and Figure 2 show the results. Fable names the exact month of 202 of the 206 blinded




Figure 2: Dating accuracy. Fable on blinded and paraphrased statements against a text-similarity baseline that has no memory. For the baseline, “exact meeting” is the upper bound described in the text.

statements (98%) and places every one within three months. The four misses (2002-05-07, 2003-09-16, 2004-03-16, 2010-12-14) are each an adjacent meeting. On the zero-lower-bound statements, where the rate level gives no clue, it is exact on 55 of 56. It reports high confidence on 205 statements. Its replies also restore details that blinding removed: for the blinded statement of 18 December 2013 it names Eric Rosengren as the dissenter, and for 25 June 2003 it says that Robert Parry dissented in favour of a larger cut. Both are correct. Paraphrasing barely helps: Fable names the exact month of 28 of the 30 paraphrased statements.

The baseline shows what the text alone can do. It places 92% of the statements within three months and 51% within one month, so dating to within a few months is not evidence of memory. Exact identification is more telling. The baseline’s upper bound on exact recovery is 43% (89 of 206), and 28 of 56 at the zero lower bound, against Fable’s 98% and 55 of 56. The probe therefore shows that Fable has knowledge sufficient to identify the individual meeting. Whether that knowledge is of the statements themselves or of a detailed policy history cannot be separated here, but the zero-lower-bound statements, which share a rate level and differ mainly in purchase amounts and wording, are the hardest case for general knowledge, and the restored dissenters’ names point to the statements themselves.

5.2 Recognition while scoring

Dating on request shows that Fable *can* identify the statements. A more direct question is whether it does so when it is only asked for a score. I search every stored scoring reply for a date in the form “<month> [day,] <year>” or “<year> statement/meeting/decision” and check whether it is the statement’s own month and year. On blinded and paraphrased text no date is in the input, so a correct date came from the model. On the original text of the five-prompt runs, 51 early statements contained their year in a navigation line (Section 2), but not the month. The counts are lower bounds: most prompts ask for the number only, and replies were stored cut at 300 characters.

Table 6 reports the counts. Across its five-prompt runs on the original text, Fable names the statement’s own month and year in 19.6% of its replies. The share depends on whether the prompt lets it talk: 58.8% for the reason-first prompt P5, 0.8% for the analyst prompt P3. On blinded

Table 6: Scoring replies that name the statement’s date, unprompted

Scores	Text	Prompt	Replies	Name a date	Own month and year
Claude Fable 5.1, five prompts	original	all	1228	245 (20.0%)	241 (19.6%)
		P1_minimal	255	21 (8.2%)	21 (8.2%)
		P2_defined	244	43 (17.6%)	43 (17.6%)
		P3_analyst	242	2 (0.8%)	2 (0.8%)
		P4_reversed	244	32 (13.1%)	32 (13.1%)
		P5_reason_first	243	147 (60.5%)	143 (58.8%)
Claude Fable 5.1, P1 re-run	original	all	206	20 (9.7%)	20 (9.7%)
Claude Fable 5.1, P1	blinded	all	206	23 (11.2%)	23 (11.2%)
Claude Fable 5.1, P1	paraphrased	all	206	7 (3.4%)	7 (3.4%)
Gemini 2.5 Flash, five prompts	original	all	1140	0 (0.0%)	0 (0.0%)
GPT-4, P1	original	all	228	0 (0.0%)	0 (0.0%)

Notes: A reply “names a date” if it contains “<month> [day,] <year>” or “<year> statement/meeting/decision”. “Own month and year”: the named date is the statement’s own. Replies stored before 3 October 2026 were cut at 300 characters, so the shares are lower bounds. The five-prompt Fable file includes repeat runs of some cells.

text, with P1, it does so in 11.2% of replies, slightly more than the 9.7% of the same prompt on the cleaned original. On paraphrased text the share falls to 3.4%. Gemini never names a date in 1140 replies, nor does GPT-4 in 228.

Table 7 gives examples. Some replies go beyond naming the meeting and bring in information that was not in the statement: later decisions (cuts “came in January 2001”; a statement was “setting up the June 2004 hike”), the market’s reaction and later commentary (“markets had hoped for more”; “widely read as teeing up December”), and the projections released alongside the statement (the “dot plot cut from 4 to 2 hikes”). This is direct evidence that the model retrieves knowledge from outside the text, including knowledge of later events, while it scores. It does not show that this knowledge moves the score, but it shows that the scores cannot be assumed to come from the text alone.

5.3 What recognition implies

The same excerpts show a second habit: Fable often judges a statement relative to the previous one or to what markets expected (“mildly dovish relative to prior expectations”). A score formed against expectations is closer to a policy surprise than a reading of the text in isolation would be, and knowing the expectations at each meeting is itself a form of recognition.

Recognition is necessary for a model to recall a statement’s market reaction, but it is not evidence that the model does so. The finding is therefore not that Fable’s scores are contaminated. It is that the tools meant to rule contamination out do not work on this corpus. Blinding and paraphrasing were meant to create an unrecognised version of each statement, so that scores on recognised and unrecognised text could be compared. With 134 of the 138 probed hold meetings dated exactly, there is no unrecognised group to compare with.

This finding complements Glasserman and Lin (2024), who remove company names from news headlines to limit look-ahead bias. Headlines about individual firms are many, short and mostly obscure. FOMC statements are few, long, and widely read and discussed; their content identifies

Table 7: Excerpts from Fable’s scoring replies

Statement	Prompt	Excerpt of the scoring reply
2013-10-30 (blinded)	P1_minimal	...This is the October 2013 FOMC statement: continued QE at \$85B/mo...
2014-01-29 (blinded)	P1_minimal	...This is the January 2014 FOMC statement: a second measured taper...
2000-12-19 (original)	P4_reversed	...inflation to economic weakness — a clear dovish pivot signaling upcoming cuts (which came in January 2001)...
2004-05-04 (original)	P2_defined	...removed at a pace that is likely to be measured” — a clear hawkish shift setting up the June 2004 hike,...
2007-12-11 (original)	P5_reason_first	...acknowledged intensifying housing correction and rising financial strains (markets had hoped for more)...
2016-03-16 (original)	P1_minimal	...dissent (George). Overall mildly dovish relative to prior expectations (dot plot cut from 4 to 2 hikes)...
2016-11-02 (original)	P5_reason_first	...continued to strengthen” — explicit signal that a hike is approaching (widely read as teeing up December)...

Notes: Verbatim excerpts from stored scoring replies (the blinded ones from P1 on blinded text). Each was asked only for a stance score. Several cite information that was not in the statement: later decisions, the market’s reaction or commentary, or the projections released with it.

them. For documents like these, anonymisation and paraphrase do not rule out look-ahead bias. Only models whose training data stop before the documents (He et al., 2025; Drinkall et al., 2024), or documents the model cannot have seen, can do that.

5.4 Where Fable’s knowledge ends

Running the same probe on the 22 blinded statements from 2024 to 2026 locates the end of Fable’s knowledge (Appendix Table A8). It dates 19 of them exactly, including every statement from the second half of 2024 and from 2025, all with high confidence. Its confidence falls in 2026: of the four statements from the first half of 2026 it dates only one with high confidence, and it misses July 2026. Its knowledge therefore extends to about early 2026. The 2024–2025 statements, often treated as a “post-cutoff” test sample, are not unseen for this model.

6 Validity: do score changes track market surprises?

6.1 Specification

If a stance score measures what the statement says about future policy, changes in the score should move with the market’s reading of the announcement. I estimate

$$y_t = \alpha + \beta \text{dScore}_t + \varepsilon_t \quad (1)$$

over statements up to December 2023, where y_t is a market outcome at statement t . Standard errors are Newey–West (Newey and West, 1987) with a Bartlett kernel and $\lfloor 4(n/100)^{2/9} \rfloor$ lags, a common rule of thumb associated with Newey and West (1994), which is 4 meetings in these samples; p -values use the Student t distribution.

The *primary specification* uses MPS as the outcome, Fable’s dScore, the policy corpus and hold

Table 8: The primary specification and its family

Outcome	Meetings	Control	n	$\hat{\beta}$	NW t	p	Holm p	Bonferroni p
MPS	all	–	198	0.049	3.02	0.003	0.034	0.034
MPS	hold	–	132	0.033	2.39	0.018	0.199	0.217
MPS	all	change_bp	198	0.036	2.27	0.025	0.245	0.294
d_{10} (bp)	hold	–	120	5.20	1.74	0.084	0.756	1.000
d_{10} (bp)	all	change_bp	173	4.39	1.52	0.130	1.000	1.000
d_{10} (bp)	all	–	173	4.27	1.44	0.151	1.000	1.000
d_2 (bp)	all	change_bp	173	4.93	1.41	0.160	1.000	1.000
d_2 (bp)	all	–	173	5.02	1.41	0.162	1.000	1.000
d_2 (bp)	hold	–	120	3.77	1.03	0.306	1.000	1.000
MPS_ORTH	all	–	189	0.014	0.79	0.431	1.000	1.000
MPS_ORTH	hold	–	125	0.011	0.69	0.490	1.000	1.000
MPS_ORTH	all	change_bp	189	0.010	0.56	0.578	1.000	1.000
MPS, 2003+ sample	hold	–	120	0.038	2.65	0.009	–	–

Notes: Fable consensus dScore, policy corpus, statements to December 2023. Each row is one regression of the outcome on dScore (and the rate change in basis points, where shown). The family is the twelve combinations of four outcomes (MPS, MPS_ORTH, d_2 , d_{10}) and three samples; Holm and Bonferroni adjust for these twelve tests only. Newey–West t with 4 lags. The primary specification is in bold. The last row repeats it on the 2003+ meetings that also have daily yields.

meetings only, with no controls. Hold meetings are the natural test: at a rate change both the score and the surprise move with the decision, so their correlation could simply reflect the decision, which no LLM is needed to read. MPS is measured in a 30-minute window and so isolates the announcement.

This specification was frozen on 3 October 2026, after earlier versions had been estimated and seen. The choice of hold meetings, of MPS as the outcome, of the five-prompt consensus over a single prompt, of whether to control for the rate change, and of the corpus definition were all made after looking at results. The corpus definition, in particular, followed an independent review of the code which found that seven non-policy notices, then coded as hold meetings, drove the earlier daily-yield results. The adjusted p -values below correct for the twelve regressions in Fable’s primary family. They do not correct for this search, so even they overstate the evidence. The main analysis run reports 225 regression coefficients with their tests (Appendix E), and the checks in this section add more.

6.2 The primary result

Table 8 reports the primary regression and its family. On 132 hold meetings from 2000-06-28 to 2023-12-13, the slope is 0.033 percentage points of MPS per unit of dScore, with a Newey–West t of 2.39 ($p = 0.018$) and an R^2 of 5.1%. The OLS and heteroskedasticity-robust t -statistics are similar (2.66 and 2.41). In economic terms the relation is small: a one-standard-deviation change in dScore (0.243) goes with 0.008 percentage points of MPS, about 0.20 of the standard deviation of MPS at hold meetings. On the 120 hold meetings from 2003 onwards, which also have daily yields, the t -statistic is 2.65.

The result does not survive a correction for multiple testing. Its Holm-adjusted p -value (Holm, 1979) over the twelve tests is 0.199 (Bonferroni 0.217). The only regression in the family that

Table 9: Hold meetings: slopes on dScore by outcome and score

Outcome	n	Fable	Gemini	Dictionary	Fable P1	P1 blinded	P1 paraphr.	GPT-4 P1
MPS	132	0.033 (2.39)	0.003 (0.27)	-0.010 (-0.99)	0.026 (2.20)	0.017 (1.43)	0.014 (0.90)	-0.006 (-0.63)
MPS_ORTH	125	0.011 (0.69)	-0.009 (-0.91)	-0.021 (-1.75)	0.008 (0.55)	-0.001 (-0.06)	-0.009 (-0.53)	-0.016 (-1.92)
t_2 (30-min, bp)	132	3.76 (2.05)	0.03 (0.03)	-1.52 (-1.39)	2.66 (1.74)	1.72 (1.12)	1.42 (0.74)	-1.11 (-1.03)
t_{10} (30-min, bp)	132	2.98 (2.10)	0.06 (0.10)	-0.54 (-0.66)	1.85 (1.75)	1.45 (1.38)	0.96 (0.68)	-0.61 (-0.94)
d_2 (daily, bp)	120	3.77 (1.03)	-2.26 (-1.29)	-0.36 (-0.17)	2.21 (0.74)	0.06 (0.02)	-0.97 (-0.26)	-2.55 (-1.88)
d_{10} (daily, bp)	120	5.20 (1.74)	-0.15 (-0.11)	1.70 (0.87)	3.18 (1.40)	1.44 (0.61)	1.10 (0.37)	-0.05 (-0.04)

Notes: Slope of each outcome on dScore with Newey–West t in parentheses; hold meetings, policy corpus, statements to December 2023. Fable and Gemini: five-prompt consensus with prompt fixed effects. Dictionary: word-count score. Fable P1: re-score with prompt P1 on the cleaned original text; P1 blinded and P1 paraphrased: the same prompt on the blinded and paraphrased versions. GPT-4: prompt P1; it lacks a few statements, so its n is up to three lower. MPS in percentage points per unit of dScore, yields in basis points.

survives is MPS on all meetings ($t = 3.02$, Holm $p = 0.034$). That regression includes the rate decisions; controlling for the size of the rate change lowers its t to 2.27.

6.3 Other outcomes and other scores

Table 9 reports every outcome for every score on hold meetings. Three facts stand out.

First, Fable’s dScore is unrelated to MPS_ORTH ($t = 0.69$), the part of the surprise that Bauer and Swanson (2023b) show is not predictable from macroeconomic and financial data released before the meeting. Its relation is with the predictable part: on the same 125 meetings, regressing MPS minus MPS_ORTH on dScore gives $t = 4.99$ (Table 13; an exploratory check). A model that knows the economic context of each meeting would pick up exactly that component, whether it reads it from the text or recalls it.

Second, the 30-minute yield changes give t -statistics of 2.05 (t_2) and 2.10 (t_{10}), close to the MPS result. They are not independent evidence: on hold meetings their correlation with MPS is 0.94 and 0.70. The daily yield changes are not significantly related to dScore (t of 1.03 for d_2 and 1.74 for d_{10}). Daily changes also contain the day’s other news, so they are a noisier outcome.

Third, no other score shows a relation. Gemini’s dScore has t -statistics between -1.29 and 0.27 on the six outcomes, the dictionary’s are of either sign and none is significant, and GPT-4’s are all negative or near zero. Fable’s own single-prompt re-score on the cleaned text gives an MPS t of 2.20; on blinded and paraphrased text the same prompt gives 1.43 and 0.90 (Section 6.7).

6.4 Robustness

Score construction and corpus. Table 10 re-estimates the hold-meeting regressions under five alternatives. The MPS t -statistic stays between 2.18 and 2.39 in every variant, including the single

Table 10: Hold-meeting t -statistics under other score constructions and corpora

Variant	MPS		t_2		t_{10}		d_2		d_{10}	
	n	t	n	t	n	t	n	t	n	t
Default: policy corpus, prompt-demeaned consensus	132	2.39	132	2.05	132	2.10	120	1.03	120	1.74
Strict corpus (also drops 2007-08-17, 2020-03-23)	130	2.18	130	2.11	130	2.48	118	1.98	118	2.73
All files (the seven notices kept)	137	2.33	137	2.23	137	2.59	126	2.32	126	2.76
Consensus over prompts common to both statements	130	2.36	130	2.02	130	2.07	118	1.00	118	1.65
Raw consensus (no prompt fixed effects)	132	2.38	132	2.04	132	2.09	120	1.01	120	1.73
One prompt only (P1_minimal)	117	2.20	117	1.84	117	1.96	105	0.41	105	1.27

Notes: Fable dScore, hold meetings, statements to December 2023, Newey–West t . Strict corpus: policy corpus without the two unscheduled policy statements (17 August 2007 and 23 March 2020). All files: also keeps the seven notices. Matched: each pair’s change is averaged over the prompts scored at both statements. Raw: mean over available prompts without prompt fixed effects. P1: prompt P1 only.

Table 11: Placebo days for the daily yield regressions

Corpus	Outcome	Actual t	Placebo days	$ t > 1.96$	$ t \geq \text{actual}$	95th pct. $ t $	Max $ t $
Policy (default)	d_2	1.03	60	7	27	2.37	2.71
	d_{10}	1.74	60	9	15	2.64	3.50
Strict	d_2	1.98	60	6	6	2.07	2.35
	d_{10}	2.73	60	6	2	2.44	3.05
All files	d_2	2.32	60	4	1	2.01	2.39
	d_{10}	2.76	60	10	3	2.57	3.19

Notes: Fable dScore, hold meetings. Each placebo replaces the statement-day yield change with the change k trading days earlier or later ($k = \pm 1, \dots, \pm 30$) and re-runs the regression. Columns give the number of placebo days whose $|t|$ exceeds 1.96 and the actual t , and the 95th percentile and maximum of the 60 placebo $|t|$ values.

prompt P1 (2.20). The daily yields do not behave this way. Under the policy corpus their t -statistics are 1.03 (d_2) and 1.74 (d_{10}). Dropping the two unscheduled policy statements raises them to 1.98 and 2.73, and keeping the non-policy notices gives 2.32 and 2.76. A result that moves by one unit of t depending on whether two statements are in the sample cannot be called established, whichever corpus one prefers. The policy corpus was chosen before the strict variant was run. The all-files corpus is not a defensible alternative: it treats liquidity and swap-line notices as hold meetings and measures the next statement’s dScore against them.

Placebo days. If the daily yield regressions had well-behaved t -statistics, the same regression on yield changes from days without a statement should rarely give $|t| > 1.96$. Table 11 shows that it often does. Under the policy corpus, 9 of 60 placebo days give $|t| > 1.96$ for d_{10} , the 95th percentile of placebo $|t|$ is 2.64, and 15 placebo days match or exceed the actual 1.74. The actual t -statistics are ordinary draws from this distribution. In the other two corpora the picture is different: the actual d_{10} t is matched by only 2 of 60 placebo days under the strict corpus and 3 under the all-files corpus, and the actual d_2 t by 6 and 1. That is the strongest evidence for a daily yield link in this paper. It appears only in corpora that differ from the policy corpus by two borderline statements, or that keep the non-policy notices. The fat tails of the placebo distributions are a warning in their own right: GPT-4’s coarse score gives a t of -3.33 for d_2 before its cutoff (Section 7), with the wrong sign, which shows how large a spurious daily-yield t can be here.

Table 12: Influence of individual meetings

Outcome	Full t	LOO min	LOO max	t after dropping 1, 2, ..., 5 meetings (greedy)
MPS	2.39	2.13	3.08	2.13 $\rightarrow$ 1.87 $\rightarrow$ 1.58 $\rightarrow$ 1.34 $\rightarrow$ 1.12 dropped in order: 2023-12-13, 2007-08-17, 2002-08-13, 2004-01-28, 2015-10-28
d_2	1.03	0.64	1.99	0.64 $\rightarrow$ 0.40 $\rightarrow$ 0.14 $\rightarrow$ -0.09 $\rightarrow$ -0.30 dropped in order: 2023-12-13, 2015-10-28, 2004-01-28, 2015-09-17, 2019-06-19
d_{10}	1.74	1.54	2.74	1.54 $\rightarrow$ 1.38 $\rightarrow$ 1.21 $\rightarrow$ 1.03 $\rightarrow$ 0.85 dropped in order: 2023-12-13, 2015-10-28, 2013-06-19, 2011-08-09, 2015-09-17
t_2	2.05	1.79	2.60	1.79 $\rightarrow$ 1.57 $\rightarrow$ 1.36 $\rightarrow$ 1.13 $\rightarrow$ 0.91 dropped in order: 2023-12-13, 2015-10-28, 2004-01-28, 2019-01-30, 2002-08-13
t_{10}	2.10	1.90	2.48	1.90 $\rightarrow$ 1.73 $\rightarrow$ 1.58 $\rightarrow$ 1.43 $\rightarrow$ 1.28 dropped in order: 2023-12-13, 2015-10-28, 2019-01-30, 2004-01-28, 2013-06-19

Notes: Fable dScore, hold meetings, policy corpus, Newey–West t . LOO: range of t when one meeting is left out. Greedy: at each step the meeting whose removal lowers t most is dropped.

Influence. Table 12 and Figure 3 show how much the MPS result depends on a few meetings. Leaving out one meeting at a time gives t -statistics between 2.13 and 3.08. Dropping the five most influential meetings one after another gives $2.39 \rightarrow 2.13 \rightarrow 1.87 \rightarrow 1.58 \rightarrow 1.34 \rightarrow 1.12$. This greedy drop removes the observations most favourable to the result, so it weakens any effect of this size. To calibrate it, I simulate 200 data sets from the fitted line with resampled residuals, so that the true slope equals the estimate, and apply the same drop. The median t -statistic falls from 2.75 to 1.15, and 49% of the simulations end at or below the actual 1.12. The influence analysis is therefore what a real effect of this size would produce, and it is not evidence against the result. The relation is also not driven by outliers: winsorising both variables at the 2.5th and 97.5th percentiles gives $t = 2.54$, the Spearman rank correlation is 0.28, and a permutation test that shuffles dScore across hold meetings gives $p = 0.006$ (which, like the other p -values, ignores the specification search).

The most influential meeting for every outcome is the statement of 13 December 2023, widely read as signalling the end of the hiking cycle. Among the others are the statement of 28 January 2004, which replaced the promise to keep rates low for a “considerable period” with a promise to be “patient”, and that of 28 October 2015, which pointed to a possible hike at the next meeting. These meetings are not recognised more often than others: Fable names their month and year unprompted in 8 of 29 scoring replies (28%), against 168 of 708 (24%) at the other hold meetings. Conversely, Fable’s dScore misses the two largest asset-purchase surprises among the hold meetings: it is -0.08 for 18 March 2009, when the 10-year yield changed by -51 basis points in a day, and 0.00 for 18 September 2013 (-17 basis points).

Subperiods and other checks. Table 13 reports exploratory checks that an independent review asked for. The relation is concentrated in the last third of the sample: the t -statistic is 1.32 for 2000–2007 ($n = 32$), 0.42 for 2008–2015 ($n = 58$) and 3.40 for 2016–2023 ($n = 42$). The score level, rather than its change, is not significantly related to MPS ($t = 1.38$). Controlling for the previous statement’s MPS or for the previous dScore leaves the t -statistic at 2.39 and 2.29, so the result does not come from persistence in either series. The concentration in 2016–2023 coincides with the period in which the Chair held a press conference after every meeting from 2019; the 30-minute window of MPS largely precedes the press conference, but the daily yield changes include it, and I do not control for it.




Figure 3: MPS against Fable’s dScore at 132 hold meetings, 2000–2023, with the fitted line of the primary regression. Filled points: the five meetings whose removal lowers the t -statistic most.

Information beyond the surprise. Does dScore tell us anything about yields that the surprise does not? It does not. Adding dScore to a regression of daily yield changes on MPS raises the R^2 by between 0.04 and 0.64 percentage points, with no $|t|$ above 1.06 (Appendix Table A3).

6.5 Fable against Gemini

Gemini’s dScore correlates with Fable’s at 0.739, yet on its own it has no relation to any market outcome. Table 14 puts both in one regression, with each dScore standardised. Fable’s coefficient has t -statistics between 2.76 and 3.41, and Gemini’s turns negative, between -3.33 and -2.66 . With two correlated regressors this joint regression is close to a regression on the difference between them: the market-related information sits in the part of Fable’s score change that Gemini does not share. The negative Gemini coefficient does not mean that Gemini’s score is informative with the wrong sign. The difference between the two models’ standardised slopes, estimated from separate regressions, has a block-bootstrap z between 2.23 and 2.59 (moving blocks of four meetings, 2,000 draws; Künsch, 1989), and every 95% interval excludes zero. Without the five meetings most influential for the primary result (Table 12), the joint t -statistics are between 2.29 and 2.75 for Fable and between -2.80 and -2.15 for Gemini, and the gap z is between 1.83 and 2.16. These comparisons were added after the primary specification was frozen and lie outside its family, so they are exploratory.

The difference between the two models is therefore suggestive, and more than a difference in significance. It has two readings. Fable may read the statements better. Or it may know more about them, including what followed. Section 5 shows that Fable recognises the statements. Gemini never names a date while scoring, but it was not given the dating probe, so whether it could recognise them is untested. The human labels below help to tell the readings apart, but only a little.

6.6 Agreement with human labels

If Fable’s edge over Gemini came from better reading, Fable should also agree better with human readers. Table 15 compares each score’s level with the mean of the human labels of Shah et al.

Table 13: Exploratory checks

Check (outcome MPS, hold meetings)	n	$\hat{\beta}$	NW t	p	Sample
Subperiod 2000–2007	32	0.036	1.32	0.196	2000-06–2007-08
Subperiod 2008–2015	58	0.012	0.42	0.673	2008-06–2015-10
Subperiod 2016–2023	42	0.049	3.40	0.002	2016-01–2023-12
Score level instead of its change	132	0.010	1.38	0.170	2000-06–2023-12
Primary + previous statement’s MPS (dScore coef.)	132	0.033	2.39	0.019	2000-06–2023-12
Primary + lagged dScore (dScore coef.)	132	0.033	2.29	0.024	2000-06–2023-12
Joint: paraphrased dScore	132	0.015	0.98	0.331	2000-06–2023-12
and clean – paraphrased	132	0.045	3.59	<0.001	2000-06–2023-12
Joint: blinded dScore	132	0.021	1.83	0.070	2000-06–2023-12
and clean – blinded	132	0.060	2.25	0.026	2000-06–2023-12
Outcome MPS – MPS_ORTH (predictable part)	125	0.022	4.99	<0.001	2000-06–2023-12
MPS on the same meetings	125	0.033	2.39	0.018	2000-06–2023-12
Primary, x and y winsorised at 2.5/97.5%	132	0.034	2.54	0.012	2000-06–2023-12

Notes: Outcome MPS (or its predictable part, MPS – MPS_ORTH), hold meetings, policy corpus, Newey–West t . All rows except the joint rows use Fable’s consensus dScore; “coef.” rows report the dScore coefficient in a regression with the stated control. Joint rows: MPS regressed in one equation on Fable’s single-prompt dScore from the paraphrased (or blinded) text and on the difference between the clean and the paraphrased (blinded) dScore. These checks were added after the primary specification was frozen and count towards the regression total.

Table 14: Fable against Gemini, hold meetings

Outcome	n	Joint: Fable		Joint: Gemini		Gap in per-SD slopes		
		$\hat{\beta}$	NW t	$\hat{\beta}$	NW t	Gap	z	95% interval
MPS	132	0.020	3.41	−0.014	−2.66	0.008	2.23	[0.002, 0.016]
d_2	120	3.11	2.76	−2.94	−3.32	1.60	2.59	[0.30, 2.70]
d_{10}	120	2.94	3.09	−2.20	−3.32	1.36	2.56	[0.38, 2.45]
t_2	132	2.42	3.21	−1.84	−2.79	1.00	2.38	[0.24, 1.87]
t_{10}	132	1.90	3.33	−1.43	−3.33	0.78	2.46	[0.27, 1.50]

Notes: Both dScores standardised on the common hold meetings. Joint: one regression of the outcome on both, Newey–West t . Gap: Fable’s minus Gemini’s per-standard-deviation slope from separate regressions; z and the 95% percentile interval from a moving-block bootstrap over meetings in time order (block length 4, 2,000 draws). MPS in percentage points, yields in basis points, per standard deviation of dScore.

(2023) on the 100 statements with at least two labelled sentences. All three LLMs agree moderately with the labels, and Fable does not agree best: its correlation is 0.459, against 0.511 for Gemini and 0.514 for GPT-4, and 0.366 against 0.516 and 0.455 on hold meetings. The dictionary reaches 0.309.

This benchmark is limited. The labels were made for sentences, not statements; only a few sentences per statement are labelled, many of them boilerplate; and the test is on levels, while the market test is on changes. The labelled sentences differ between neighbouring statements, so a change test with these labels mostly measures which sentences happen to be labelled (all models’ correlations with the change in the mean label are between 0.048 and 0.081). Within these limits, the human benchmark does not single out Fable as the better reader. Its market edge is not matched by an edge in agreement with humans, which fits “knows more” at least as well as “reads better”.

Table 15: Correlation of score levels with published human sentence labels

Score	All statements		Hold meetings	
	n	r	n	r
Claude Fable 5.1 (five prompts)	100	0.459	76	0.366
Gemini 2.5 Flash (five prompts)	100	0.511	76	0.516
Dictionary word count	100	0.309	76	0.199
GPT-4 (P1)	97	0.514	73	0.455

Notes: Statements with at least two sentences that appear word for word among the hand-labelled sentences of Shah et al. (2023) (hawkish +1, neutral 0, dovish -1); r is the correlation of the model’s score level with the mean label, 2004–2022. GPT-4 lacks scores for three statements.

6.7 Blinded and paraphrased text

It is tempting to read the fall in the MPS t -statistic from 2.20 on the clean text to 1.43 on blinded and 0.90 on paraphrased text as evidence that recalled information was removed. The evidence is mixed. The falls are not clearly significant: of the 6 bootstrap gaps between the clean slope and the blinded or paraphrased one (three outcomes, two text versions), with z -statistics between 0.97 and 2.15, only 1 has $|z|$ above 1.96 (the d_2 gap for blinding), and every 95% percentile interval includes zero (Appendix Table A4). Recognition on request hardly changes: Fable still dates 98% of blinded and 93% of paraphrased statements. But recognition while scoring does fall: Fable names the statement’s month and year in 9.7% of its P1 replies on the clean text, 11.2% on blinded and 3.4% on paraphrased text. And paraphrasing removes legitimate information: markets react to exact wording such as “considerable period” or “patient”, which a paraphrase must change.

The last point can be made concrete. Splitting the clean dScore into the paraphrased dScore and the difference between the two, and entering both in one regression for MPS, gives a t -statistic of 0.98 for the paraphrased part and 3.59 for the difference (Table 13); for blinding the figures are 1.83 and 2.25. The market-related part of the score is concentrated in what paraphrasing and blinding change. For the paraphrase, that fits both readings: exact wording carries information, and identity cues may trigger recall. Blinding is more telling, because it leaves the wording of the policy paragraphs intact and removes mainly dates, names, vote rosters and named events. That the part of the score blinding removes is related to MPS ($t = 2.25$) points towards identity cues, or towards the context they carry, rather than towards wording. This is the strongest indication of recall in the paper. It is weak evidence: the blinded and clean dScores correlate at 0.913, the check is exploratory, and the gap between the clean and blinded slopes is not significant on its own.

6.8 Summary

Fable’s dScore has a positive and small relation to the policy surprise at hold meetings. It changes little with how the score is built or how the corpus is defined, it is not driven by outliers, and its sensitivity to influential meetings is what a real effect of this size would show. But it does not survive a multiple-testing correction, it was found after a specification search, and it is concentrated in 2016–2023. The daily yield link depends on the corpus definition. Fable outperforms Gemini on market data but not on human labels. None of this separates reading from remembering, and the one check that leans either way (the part of the score that blinding removes) leans towards

remembering.

7 What can and cannot be tested

The central question, whether Fable’s market link comes from reading or from remembering, has a clean answer only if some statements are unknown to the model. This section goes through the ways to get such statements and explains why none of them works with the data at hand.

7.1 Waiting for new statements

The policy corpus has 21 statement pairs from 2024 to July 2026, 15 of them at hold meetings. The surprise series ends in 2023, so only the daily yields can be used. No score shows a significant relation (Appendix Table A6): on hold meetings Fable’s t -statistics are 1.06 (d_2) and 0.76 (d_{10}), Gemini’s 1.17 and 1.03.

This is not evidence against the pre-2024 result, because the test has almost no power. If Fable’s pre-2024 R^2 were the true effect, the current post-2024 sample would detect it with a probability of 8% to 11% (Appendix Table A7). Reaching 80% power would take between 280 and 480 meetings, or 37 to 56 more years at the current pace. These R^2 values come from the sample that first suggested the effect and are, if anything, too high, so the true requirement is likely larger. Worse, Section 5.4 showed that Fable knows the 2024–2025 statements. At most the four statements from March to July 2026 may be unseen.

7.2 A model with a documented training cutoff

GPT-4 is listed with a training cutoff of September 2021. In principle it allows a comparison within one model: statements before the cutoff it may have memorised, statements after it cannot have. In practice it fails as an instrument. Its 225 scores take only 10 distinct values; 98 of them (44%) are exactly -1 ; 67% of its hold-meeting dScores are exactly zero (Fable: 5 of 132); and 40 of the 56 zero-lower-bound statements sit at the floor of -1 . A score that rarely changes between meetings cannot track meeting-to-meeting surprises.

Its regressions show no positive relation before the cutoff (hold-meeting t of -0.24 for MPS and -3.33 for d_2 , the latter with the wrong sign) or after it ($|t| \leq 1.33$ on 39 or fewer statements; Appendix Table A5). On the same post-cutoff statements Fable’s t -statistics are between 1.77 and 2.37, but Fable may know those statements. GPT-4’s null result therefore says nothing about whether having seen the statements is enough to produce a market link. It shows only that GPT-4 with this prompt is too coarse to be a control.

7.3 Recognised against unrecognised statements

The planned test was to compare the market link on statements that Fable recognises with the link on statements it does not. Section 5 showed that the second group is nearly empty (4 meetings): Fable dates 134 of the 138 probed hold meetings exactly, and blinding and paraphrasing do not change that. Restricting the regressions to the recognised statements simply repeats the main result.

7.4 What would settle the question

Three designs could separate reading from remembering. First, chronologically consistent models (He et al., 2025; Drinkall et al., 2024) are trained only on text available before each date and so cannot remember what followed a statement. Scoring the corpus with such models, using the same prompts, would give an uncontaminated version of the market test. Whether such models produce scores fine enough to detect an effect of this size is an open question; GPT-4’s coarse output shows that it cannot be taken for granted. Second, a second frontier model with a different training cutoff would show whether the link appears only for statements a model knows. A dating probe on such a model is cheap (a few dollars) but was not run here. Third, human labels of meeting-to-meeting *changes* would provide a validity anchor that neither markets nor model memory can touch. A sheet of 40 randomly drawn hold-meeting pairs is prepared in the repository but has not been labelled.

8 Conclusion

Can an applied researcher trust an LLM stance score of FOMC statements, and can they tell whether the model read the text or remembered it? For the frontier model studied here, the answer to the first question is largely yes, as a matter of measurement: its scores barely depend on the wording of the prompt, and they survive the removal of the rate decision. That consistency may itself owe something to recognition, and the few statements the model probably has not seen are too few to tell. The answer to the second question is no, at least with the tools tested here. The model identifies almost every statement even after blinding and paraphrasing, and it sometimes cites later events while scoring. Its scores are weakly related to monetary policy surprises at no-change meetings, but that relation does not survive a multiple-testing correction, is concentrated in recent years, and cannot be attributed to reading rather than remembering.

Four practical lessons follow for work that uses LLMs to measure text.

1. Report reliability where it matters. Agreement on raw levels can be driven by the obvious content (here, the rate decision). Report agreement after removing it and on changes, and report variation across models as well as prompts: for changes, the model mattered about twice as much as the prompt.
2. Test recognition directly. A dating probe calibrated against a text-similarity baseline, and a search of the scoring replies for unprompted identification, are cheap and should accompany any LLM measure of well-known documents.
3. Do not rely on anonymisation or paraphrase to rule out look-ahead bias for well-known documents. They do not prevent identification. Only models trained on data that stop before the documents, or documents the model cannot have seen, can rule it out.
4. Treat a frontier model’s market validation on well-known documents as possibly contaminated until such a test confirms it, and report the specification search, the placebo distribution of the test statistic and the influence of individual observations, calibrated against what a real effect would show.

Limitations. The study uses one frontier model; another may behave differently, and the dating probe was not run on a second frontier model. The crossed reliability block covers 40 statements from 2015–2019. Replies were stored cut at 300 characters until October 2026, so the counts of unprompted recognition are lower bounds and some long replies could not be re-parsed. The two latest statements could not be checked against the Federal Reserve’s website. The paper uses no fine-tuned text classifier (such as FinBERT) as a further baseline, does not split the surprises into the forward-guidance and asset-purchase factors of Swanson (2021), and has no human labels of meeting-to-meeting changes. It does not control for the Chair’s press conferences, which affect the daily yield changes, and it does not compare the scores with the text measures of Lucca and Trebbi (2009), Hansen and McMahon (2016) or Gardner et al. (2022). Finally, the models were accessed through a commercial interface in 2026; providers can change the model behind a name, so exact replication of the scores may not be possible later, although every number in this paper can be recomputed from the stored scores.

References

- Mikael Apel and Marianna Blix Grimaldi. The information content of central bank minutes. Sveriges Riksbank Working Paper Series 261, Sveriges Riksbank, April 2012. Published as “How Informative Are Central Bank Minutes?”, *Review of Economics* 65(1), 53–76, 2014.
- Michael D. Bauer and Eric T. Swanson. An alternative explanation for the “Fed information effect”. *American Economic Review*, 113(3):664–700, 2023a. doi: 10.1257/aer.20201220.
- Michael D. Bauer and Eric T. Swanson. A reassessment of monetary policy surprises and high-frequency identification. *NBER Macroeconomics Annual*, 37:87–155, 2023b. doi: 10.1086/723574. NBER Macroeconomics Annual 2022.
- Nicholas Carlini, Florian Tramèr, Eric Wallace, Matthew Jagielski, Ariel Herbert-Voss, Katherine Lee, Adam Roberts, Tom Brown, Dawn Song, Úlfar Erlingsson, Alina Oprea, and Colin Raffel. Extracting training data from large language models. In *30th USENIX Security Symposium (USENIX Security 21)*, pages 2633–2650. USENIX Association, 2021.
- Felix Drinkall, Eghbal Rahimikia, Janet B. Pierrehumbert, and Stefan Zohren. Time Machine GPT. In Kevin Duh, Helena Gomez, and Steven Bethard, editors, *Findings of the Association for Computational Linguistics: NAACL 2024*, pages 3281–3292, Mexico City, Mexico, 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.findings-naacl.208.
- Ben Gardner, Chiara Scotti, and Clara Vega. Words speak as loudly as actions: Central bank communication and the response of equity prices to macroeconomic announcements. *Journal of Econometrics*, 231(2):387–409, 2022.
- Paul Glasserman and Caden Lin. Assessing look-ahead bias in stock return predictions generated by GPT sentiment analysis. *The Journal of Financial Data Science*, 6(1):25–42, 2024. doi: 10.3905/jfds.2023.1.143.

- Shahriar Golchin and Mihai Surdeanu. Time travel in LLMs: Tracing data contamination in large language models. In *The Twelfth International Conference on Learning Representations (ICLR 2024)*, 2024. arXiv:2308.08493.
- Refet S. Gürkaynak, Brian Sack, and Eric T. Swanson. Do actions speak louder than words? The response of asset prices to monetary policy actions and statements. *International Journal of Central Banking*, 1(1):55–93, 2005.
- Anne Lundgaard Hansen and Sophia Kazinnik. Can ChatGPT decipher FedSpeak? Working paper, SSRN 4399406, 2023. URL https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4399406.
- Stephen Hansen and Michael McMahon. Shocking language: Understanding the macroeconomic effects of central bank communication. *Journal of International Economics*, 99(S1):S114–S133, 2016. doi: 10.1016/j.jinteco.2015.12.008.
- Andrew F. Hayes and Klaus Krippendorff. Answering the call for a standard reliability measure for coding data. *Communication Methods and Measures*, 1(1):77–89, 2007. doi: 10.1080/19312450709336664.
- Songrun He, Linying Lv, Asaf Manela, and Jimmy Wu. Chronologically consistent large language models. Working paper, arXiv:2502.21206, 2025.
- Sture Holm. A simple sequentially rejective multiple test procedure. *Scandinavian Journal of Statistics*, 6(2):65–70, 1979. URL <https://www.jstor.org/stable/4615733>.
- Marek Jarociński and Peter Karadi. Deconstructing monetary policy surprises—The role of information shocks. *American Economic Journal: Macroeconomics*, 12(2):1–43, 2020. doi: 10.1257/mac.20180090.
- Klaus Krippendorff. *Content Analysis: An Introduction to Its Methodology*. SAGE Publications, 4th edition, 2018. ISBN 9781506395661.
- Hans R. Künsch. The jackknife and the bootstrap for general stationary observations. *The Annals of Statistics*, 17(3):1217–1241, 1989. doi: 10.1214/aos/1176347265.
- Alejandro Lopez-Lira and Yuehua Tang. Can ChatGPT forecast stock price movements? Return predictability and large language models. Working paper, arXiv:2304.07619, 2023.
- David O. Lucca and Francesco Trebbi. Measuring central bank communication: An automated approach with application to FOMC statements. NBER Working Paper 15367, National Bureau of Economic Research, September 2009.
- Jens Ludwig, Sendhil Mullainathan, and Ashesh Rambachan. Large language models: An applied econometric framework. *Annual Review of Economics*, 18:283–316, 2026. doi: 10.1146/annurev-economics-120925-105620. Earlier version: NBER Working Paper 33344 (2025).
- Emi Nakamura and Jón Steinsson. High-frequency identification of monetary non-neutrality: The information effect. *Quarterly Journal of Economics*, 133(3):1283–1330, 2018. doi: 10.1093/qje/qjy004.

- Whitney K. Newey and Kenneth D. West. A simple, positive semi-definite, heteroskedasticity and autocorrelation consistent covariance matrix. *Econometrica*, 55(3):703–708, 1987. doi: 10.2307/1913610.
- Whitney K. Newey and Kenneth D. West. Automatic lag selection in covariance matrix estimation. *Review of Economic Studies*, 61(4):631–653, 1994. doi: 10.2307/2297912.
- Suproteem K. Sarkar and Keyon Vafa. Lookahead bias in pretrained language models. Working paper, SSRN 4754678, 2024. URL https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4754678.
- Agam Shah, Suvan Paturi, and Sudheer Chava. Trillion Dollar Words: A new financial dataset, task & market analysis. In Anna Rogers, Jordan Boyd-Graber, and Naoaki Okazaki, editors, *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 6664–6679, Toronto, Canada, 2023. Association for Computational Linguistics. doi: 10.18653/v1/2023.acl-long.368.
- Patrick E. Shrout and Joseph L. Fleiss. Intraclass correlations: Uses in assessing rater reliability. *Psychological Bulletin*, 86(2):420–428, 1979. doi: 10.1037/0033-2909.86.2.420.
- Eric T. Swanson. Measuring the effects of Federal Reserve forward guidance and asset purchases on financial markets. *Journal of Monetary Economics*, 118:32–53, 2021.

Appendix

A Prompts

A.1 Scoring prompts

Each prompt is sent with the statement text in place of `{statement}`, at temperature 0. The reply of P4 is multiplied by -1 . P5 is given a larger reply budget so that it can write its reasons.

P1_minimal

Score the monetary policy stance of the following FOMC statement on a scale from -1 to $+1$, where -1 is dovish and $+1$ is hawkish.

Reply with only the number.

STATEMENT:
{statement}

P2_defined

Below is a statement from the Federal Open Market Committee.

A hawkish stance leans toward tighter policy and higher interest rates, typically
↪ motivated by concern about inflation. A dovish stance leans toward looser policy and
↪ lower interest rates, typically motivated by concern about growth and employment.

Rate the stance on a continuous scale from -1 (maximally dovish) to $+1$ (maximally
↪ hawkish). Reply with only the number.

STATEMENT:
{statement}

P3_analyst

You are a fixed income analyst. Read the FOMC statement below and tell me how hawkish it
↪ is, as a single number between -1 (dovish) and $+1$ (hawkish).

Respond with the number and nothing else.

STATEMENT:
{statement}

P4_reversed (answer sign flipped back in code)

Assess the policy stance expressed in this FOMC statement.

Use a scale from -1 to +1 where +1 means strongly DOVISH and -1 means strongly HAWKISH.

Output only the numeric value.

STATEMENT:
{statement}

P5_reason_first

Read the FOMC statement below. Briefly note the main hawkish elements and the main dovish
↔ elements. Then give an overall stance score from -1 (dovish) to +1 (hawkish).

End your reply with a line in exactly this format:
SCORE: <number>

STATEMENT:
{statement}

A.2 Dating probe

For the paraphrased statements, “{reworded}” is replaced by “(it may have been reworded)”; for blinded statements it is empty.

Below is a statement released by the Federal Open Market Committee after one of its meetings{reworded}. Month names, years, member names and named events have been replaced by placeholders such as [month], [year+1] or [event]; [year+N] means N years after the year of the meeting itself.

Using anything you know, identify the meeting at which this statement was released. Give your single best guess even if unsure.

End your reply with exactly one line in this format:
ANSWER: YEAR=YYYY MONTH=MM CONFIDENCE=low|medium|high

Statement:
{statement}

A.3 Paraphrase prompt (Gemini 2.5 Flash)

{decision} states the true change in the federal funds target, for example “the federal funds rate target was left unchanged at this meeting.”

Rewrite the following central bank policy statement so that a reader who knows the original by heart could not recognise it, while an economist would draw exactly the same conclusions from it.

Rules:
- Keep ALL of the economic content: the policy decision, the direction and

- size of any change, the assessment of growth, the labour market, inflation and financial conditions, the balance of risks, any guidance about future policy (including numerical thresholds in that guidance), and any dissents with the direction the dissenter preferred.
- Change the structure: start with the decision, then the economic assessment, then the outlook and risks, then guidance, then dissents. Write in plain prose in your own words, as different central bank staff would. Do not reuse any phrase of four or more words from the original, and do not follow its sentence order.
 - Do NOT state absolute levels of the policy interest rate, its target range, the discount rate, or the size of asset holdings or purchase programmes. State them only relative to the previous meeting: "raised its policy rate by a quarter of a percentage point", "kept its policy rate unchanged", "cut the monthly pace of its Treasury and mortgage bond purchases by \$5 billion each". The 2 percent inflation objective may be kept.
 - Keep placeholders such as [month], [year+1], [member], [event] as they are, and do not add any dates, names or events.
 - Add NOTHING that is not in the original. If the original says nothing about the labour market, inflation, risks, guidance or dissents, neither do you; do not mention a 2 percent objective unless the original does. A short original gives a short rewrite.
 - Do not add interpretation, commentary or a hawkish/dovish label; do not make the tone stronger or weaker than the original.
 - About two thirds of the original length. Output only the rewritten text.

Fact to use for the decision (from the official record, do not contradict it): {decision}

Statement:
{statement}

B Corpus classification

Table A1: Statements that are not scheduled post-meeting statements

Date	Type	Note
2001-01-03	unscheduled rate	conference call; -50 bp
2001-04-18	unscheduled rate	conference call; -50 bp
2001-09-17	unscheduled rate	after 11 September 2001; -50 bp; no MPS in the Bauer–Swanson data
2007-08-10	notice	liquidity notice (reserves through open market operations); no stance change
2007-08-17	unscheduled policy	FOMC stance statement: downside risks ‘increased appreciably’; no funds-rate change
2008-01-22	unscheduled rate	-75 bp
2008-03-11	notice	Term Securities Lending Facility and swap lines; liquidity measures
2008-10-08	unscheduled rate	coordinated central bank cut; -50 bp
2010-05-09	notice	dollar swap lines re-established; a Sunday
2019-10-11	notice	technical bill purchases; text says they ‘do not represent a change in the stance’
2020-03-03	unscheduled rate	-50 bp
2020-03-15	unscheduled rate	-100 bp plus purchases; a Sunday (yields from Monday 2020-03-16)
2020-03-23	unscheduled policy	open-ended Treasury/MBS purchases; no funds-rate change
2020-03-31	notice	repo facility for foreign and international monetary authorities facility
2020-08-27	notice	framework: Statement on Longer-Run Goals and Strategy
2025-08-22	notice	framework: Statement on Longer-Run Goals and Strategy

Notes: All other 212 files are scheduled post-meeting statements. The policy corpus keeps every type except “notice”; the strict corpus also drops “unscheduled policy”. Classification by hand after reading each text.

C Parse audit

The scorers originally took the last number in a reply and clipped it to $[-1, 1]$. Since 3 October 2026 the parser requires the reply to end in the number (or to contain a “SCORE:” line) and rejects truncated numbers and values outside the range. Every stored reply shorter than the old 300-character storage limit was re-parsed with the strict rule; cells whose stored score differs from the re-parsed one are dropped. Replies stored at the 300-character limit cannot be re-parsed: 278 of Fable’s 1228 stored replies and 228 of Gemini’s 1140. One of them was found wrong by reading it: the reply was cut off after a rate level (“1.75%”), which the old parser read as 1 and flipped to -1 . Among the other cut-off replies, the scores of exactly ± 1 , where a clipped rate level would show up, agree in sign with the mean of the statement’s other prompts in 4 of 4 cases for Fable and 19 of 19 for Gemini.

Table A2: Cells dropped by the parse audit

File	Statement	Prompt	Stored	Reply (start)
scores_fable.csv	2000-11-15	P4_reversed	0.00	-0.
scores_fable.csv	2002-03-19	P4_reversed	-1.00	This is the March 2002 FOMC statem...
gpt4_p1.csv	2017-05-03	P1_minimal	0.00	0.
gpt4_p1.csv	2017-09-20	P1_minimal	0.00	0.
gpt4_p1.csv	2017-11-01	P1_minimal	0.00	0.

D Additional results

Table A3: Daily yields on MPS, with and without dScore

Outcome	Meetings	Score	n	R^2 base	R^2 + dScore	Gain (pp)	$\hat{\beta}$	NW t
d_2	all	fable	173	0.337	0.339	0.18	1.40	0.51
d_2	all	gemini	173	0.337	0.340	0.32	-1.50	-0.88
d_2	all	clean	173	0.337	0.337	0.05	0.71	0.30
d_2	all	blind	173	0.337	0.337	0.01	-0.34	-0.16
d_2	all	para	173	0.337	0.337	0.04	-0.73	-0.26
d_2	hold	fable	120	0.505	0.506	0.04	-0.52	-0.21
d_2	hold	gemini	120	0.505	0.516	1.05	-2.18	-1.26
d_2	hold	clean	120	0.505	0.508	0.21	-1.03	-0.50
d_2	hold	blind	120	0.505	0.512	0.66	-1.88	-1.03
d_2	hold	para	120	0.505	0.513	0.80	-2.56	-0.97
d_{10}	all	fable	173	0.134	0.138	0.38	2.12	0.84
d_{10}	all	gemini	173	0.134	0.136	0.19	-1.22	-0.77
d_{10}	all	clean	173	0.134	0.136	0.15	1.25	0.58
d_{10}	all	blind	173	0.134	0.134	0.01	0.38	0.18
d_{10}	all	para	173	0.134	0.134	0.00	-0.06	-0.02
d_{10}	hold	fable	120	0.127	0.133	0.64	2.59	1.06
d_{10}	hold	gemini	120	0.127	0.127	0.00	-0.10	-0.06
d_{10}	hold	clean	120	0.127	0.128	0.16	1.15	0.57
d_{10}	hold	blind	120	0.127	0.127	0.00	0.20	0.10
d_{10}	hold	para	120	0.127	0.127	0.00	0.07	0.03

Notes: Statements to December 2023, policy corpus. “All” controls for MPS and the rate change; “hold” for MPS. Gain = R^2 with dScore minus R^2 without, in percentage points. Newey–West t for the dScore coefficient.

Table A4: Clean minus blinded or paraphrased slopes, hold meetings

Outcome	Other score	n	$\hat{\beta}$ clean	$\hat{\beta}$ other	Gap	z	95% interval
MPS	blinded	132	0.026	0.017	0.008	1.45	[-0.003, 0.020]
MPS	paraphrased	132	0.026	0.014	0.012	1.26	[-0.007, 0.030]
d_2	blinded	120	2.21	0.06	2.15	2.15	[-0.10, 3.85]
d_2	paraphrased	120	2.21	-0.96	3.17	1.70	[-0.48, 7.04]
d_{10}	blinded	120	3.18	1.44	1.74	1.54	[-0.46, 4.17]
d_{10}	paraphrased	120	3.18	1.10	2.08	0.97	[-2.24, 6.41]

Notes: Fable, prompt P1. Gap = slope on the clean-text dScore minus slope on the blinded or paraphrased dScore, same meetings. Standard error, z and percentile interval from a moving-block bootstrap over meetings in time order (block length 4, 2,000 draws).

Table A5: GPT-4 before and after its listed training cutoff (September 2021)

Outcome	Window	GPT-4 (P1)			Fable (consensus)	
		n	$\hat{\beta}$	t	$\hat{\beta}$	t
MPS	all, pre	175	0.005	0.57		
MPS	hold, pre	122	-0.002	-0.24		
MPS	all, post	18	-0.010	-0.31	0.119	2.16
d_2	all, pre	150	-2.62	-1.79		
d_2	hold, pre	110	-3.56	-3.33		
d_2	all, post	39	-0.25	-0.09	11.10	1.90
d_2	hold, post	22	2.05	0.56	12.95	1.90
d_{10}	all, pre	150	-1.54	-1.12		
d_{10}	hold, pre	110	-1.08	-0.95		
d_{10}	all, post	39	-0.35	-0.18	6.27	1.77
d_{10}	hold, post	22	2.65	1.33	8.85	2.37

Notes: Pre: statements up to 30 September 2021; post: after it. MPS ends in December 2023, so the post-cutoff MPS sample is short. Post-cutoff rows use heteroskedasticity-robust t (no lags, few observations). Fable is shown on the same post-cutoff statements; it may know them.

Table A6: Statements from 2024 onwards (daily yields only)

Outcome	Meetings	n	Fable t	Gemini t	GPT-4 t
d_2	all	21	0.78	0.74	-0.43
d_2	hold	15	1.06	1.17	0.12
d_{10}	all	21	0.12	-0.40	-0.85
d_{10}	hold	15	0.76	1.03	0.04

Notes: Policy corpus, 21 statement pairs from January 2024 to July 2026. Heteroskedasticity-robust t .

Table A7: Power of the post-2024 test

Outcome	Meetings	Pre-2024 R^2	n for 80% power	n now	Power now	Pairs per year	Years more
d_2	all	0.024	321	21	11%	8.1	37
d_2	hold	0.025	314	15	9%	5.8	51
d_{10}	all	0.016	480	21	8%	8.1	56
d_{10}	hold	0.028	280	15	10%	5.8	46

Notes: Power of a two-sided 5% test if Fable’s pre-2024 R^2 were the true effect (normal approximation). “Pairs per year”: post-2024 statement pairs per year so far.

Table A8: Dating probe on the 2024–2026 statements (blinded)

Half-year	n	Exact month	Confidence “high”	Misses (error in months, confidence)
2024H1	4	3	3	2024-03-20 (-2, medium)
2024H2	4	4	4	—
2025H1	4	4	4	—
2025H2	5	5	5	—
2026H1	4	3	1	2026-04-29 (-1, medium)
2026H2	1	0	0	2026-07-29 (-1, low)

Table A9: Five-prompt scores of the seven latest statements

Statement	Dating probe	P1	P2	P3	P4	P5	SD	Percentile
2025-10-29	known	-0.60	-0.50	-0.50	-0.50	-0.40	0.063	0.70
2025-12-10	known	-0.50	-0.40	-0.40	-0.50	-0.40	0.049	0.62
2026-01-28	known	-0.10	-0.10	0.25	-0.15	0.20	0.169	0.94
2026-03-18	not known	0.10	-0.10	0.10	-0.10	0.15	0.108	0.89
2026-04-29	not known	-0.10	-0.10	0.20	-0.10	0.10	0.126	0.91
2026-06-17	not known	0.40	0.45	0.50	0.40	0.40	0.040	0.44
2026-07-29	not known	0.40	0.50	0.60	0.50	0.45	0.066	0.72

Notes: Fable, original text. “Known” = dated to the exact month with high confidence by the probe in Table A8. SD = standard deviation across the five prompts. Percentile = share of the 171 pre-2024 statements with all five prompts whose spread is at or below this one.

E Disclosures

1. **Specification search.** The primary specification was frozen on 3 October 2026, after earlier specifications had been estimated. The hold-meeting sample, the outcome, the five-prompt consensus, the treatment of the rate-change control and the corpus definition were chosen after seeing results. The default run of the analysis reports 225 regression coefficients with their tests (its printout counts 224, before the primary regression on the 2003+ sample is added); the exploratory checks of Table 13 add 13, and the Fable–Gemini comparison without influential meetings ten more; the five variant runs of Table 10 repeat the full set; and the placebo-day and influence diagnostics run several hundred more. Earlier ad-hoc checks during the project and the exploratory tests of an independent review are not counted.
2. **Corpus definition.** The policy corpus was defined after an independent review found that seven non-policy notices, then treated as hold meetings and as the base of the next statement’s dScore, drove the earlier daily-yield results. It was fixed before the strict variant was run. Under the strict corpus the daily yield t -statistics are 1.98 (d_2) and 2.73 (d_{10}), against 1.03 and 1.74 under the policy corpus.
3. **Date correction.** The statement of 27–28 June 2007 was first filed as 18 June 2007 because of the date in its web address. It is dated 28 June here.
4. **Text seen by the five-prompt runs.** For 51 statements from 2000–2005, the text scored in the five-prompt runs ended with a page-navigation line that included a year (the statement’s own year in 50 cases). Later runs used cleaned text.
5. **Paraphrases.** The paraphrase prompt was given the true change in the federal funds target. 11 paraphrases keep the 6.5 percent unemployment threshold of 2012–2014, and 2 still carry a number flagged by the automatic checks.
6. **Unverified texts.** The statements of 17 June and 29 July 2026 could not be checked against the Federal Reserve’s website. They enter the post-2024 checks, the dating probe on recent statements and the prompt-spread comparison of Section 4.
7. **Yield data.** Daily yields are FRED DGS2 and DGS10. The file matches a FRED download on all 1229 overlapping days (2021-10-06 to 2026-09-08); earlier dates were not checked against a download.
8. **Stored replies.** Until 3 October 2026 the scorers stored the first 300 characters of each reply (278 Fable and 228 Gemini replies are cut). The parse audit (Appendix C) dropped 5 cells.
9. **Missing prompts.** Fable has all five prompts for 174 policy statements and one to four for 47; the consensus uses prompt fixed effects (Table 10 shows the alternatives).
10. **Reliability residuals.** “Resid” regresses on a linear $-1/0/+1$ decision code; separate hike and cut dummies give drops of 26% to 39% for the inexpensive models and 3% for Fable.
11. **Models and access.** All models were called through OpenRouter at temperature 0. GPT-4’s training cutoff is the one OpenRouter lists. Fable’s effective knowledge end is estimated from the dating probe only.

12. **Use of AI tools.** The code, the analysis and the drafting of this paper were done with the help of an AI assistant (Claude). Every number is produced by the scripts in the repository.

F Reproduction

All results are produced, without any API call, from the stored score files:

```
sh paper/run_all.sh
```

runs the full analysis (`analysis.py`) under the default settings and the five variants, the exploratory checks (`paper/extra_checks.py`) and the table builder (`paper/make_tables.py`), then compiles this document. It takes about two minutes. `make_tables.py` writes every table and figure and a file of named numbers that the text uses, so that no number in the paper is typed by hand.